Hipstraw Faculty Harness — Evolutionary Architecture

Oct 28, 2026 · @Ashish

Hipstraw's G0 product is **Continuous Campaign Strategy & Readiness Testing**: testing a launch's campaign strategy and execution readiness against Day Zero, before and during the launch, and keeping a reviewed, versioned Position of both. The harness is one system that evolves from G0 to G4 under evidence gates. Launch tactics (Moves and Lines) are named and recorded from day one, the compiler waits for G3, and a Position is never a forecast that the campaign will succeed.

Status, scope and what changed

This is an architecture-level design for review, not an implementation specification. It recasts **Hipstraw — Faculty Harness Architecture & Evolutionary Implementation Strategy v0.1** ("HS") in the structure of the Zimplii Faculty Harness — Evolutionary Architecture v2 ("ZH"). Where this document is silent, HS applies.

What this adds to HS v0.1:

Area	HS v0.1	This architecture
Product framing	G0 promise: continuous testing of campaign strategy and readiness (§1.1)	Same, with a written scope statement and permanent exclusions (section 1)
Fixed core	Record families as a contract direction (§6)	Ten named fixed-core elements, C1–C10, built correctly once at G0 (section 4)
Launch tactics	Lines and Moves may stay declarative labels at G0; ADR-04 open	Versioned Moves with effect classes and Lines run as ordered stages from G0; every Test records its Line and deviations (section 5)
Tactic record	Not stated	C10 records Line chosen, Moves executed and every deviation, so G3 inherits an evidence base
Compiler and adaptive selection	G0 non-goal; arrives with G3 Adaptive Launch Methods	Still G3, behind the G2 → G3 gate; inherits G0 Line definitions

Area	HS v0.1	This architecture
Roles	Desk and Agent responsibilities; acceptance model open (ADR-02)	GTM leader is Sponsor/Approver; a Market Manager agent prepares and runs the Market Review (section 9)
Reports	Not stated	Rendered by the shared Report Generation Faculty
Decisions	ADR-01 to ADR-10	HR01–HR12 added and mapped to the ADRs (section 13)

Sources: HS = the uploaded `Hipstraw_Faculty_Harness_Architecture_Evolutionary_Implementation_Strategy_v0_1.docx` (28 October 2026); LS = Hipstraw Evolutionary Launch Intelligence & Orchestration Strategy v0.3 and FS = Faculty Systems Evolutionary Production Experimentation and Phased Platform Strategy v1.0, both as cited in HS; ZH = the Zimplii v2 architecture, used for structure only; Hipstraw product decisions made earlier by the product owner.

Labels: [U] decided by the product owner · **[S]** source document · **[P]** proposed here · **[A]** assumption · **[O]** open decision (section 14).

1. Product scope: Continuous Campaign Strategy & Readiness Testing

The G0 product does three things: tests a launch's campaign strategy, tests its execution readiness, and keeps a reviewed, versioned Position of both against Day Zero [S: HS §1.1, §9]. Everything in this architecture serves those three. Anything that runs campaigns, replaces the customer's channels or forecasts success is out [S: HS §1.2].

Scope statement

Continuous Campaign Strategy & Readiness Testing tests a launch's campaign strategy and execution readiness before and during the launch, measured against Day Zero. It reviews versioned, evidence-linked Positions, runs named launch Lines made of typed Moves, proposes bounded Test Plans, coordinates only authorized tests and actions, and records observed Deltas. It does not run campaigns, replace channels, analytics, CRM or ad platforms, forecast campaign success, or publish, spend or contact third parties without specific authority. [P, from HS §1]

Every Review and report carries this statement, because "launch intelligence" is easily read as campaign automation or a success forecast [P].

In and out of scope

In scope	Out of scope
Launch mandate and Day Zero intake [S: HS §5.1]	Replacing campaign channels, analytics, CRM, email or ad platforms [S: HS §1.2]
Evidence capture with provenance, recency and uncertainty [S: HS §1.3]	Publishing, spending, contacting third parties or changing budgets without specific authority [S: HS §1.2]
Reviews that separate strategy validity from operational readiness [S: HS §5.1]	Success forecasts, or readiness scores and simulations presented as proof [S: HS §1.2]
Bounded Test Plans with scoped human approval [S: HS §5.1]	A full Faculty, Facility, Desk and Agent service taxonomy at G0 [S: HS §1.2]
Coordinating authorized tests and Actions through scoped adapters [S: HS §2.1]	A playbook compiler, agent federation or method-learning engine before G3 [S: HS §1.2, §9]
Live-campaign checkpoint Reviews at T0, T+7 or configured events [S: HS §5.2]	Always-on LLM agents as a requirement at any generation [S: HS §9]
Launch tactics: Moves and Lines; later adaptive selection [P]	Secondary training use of customer data without express rights [S: HS §8]

What a Position is, and is not

A **Position** is the accepted, versioned snapshot of decision-relevant launch, market or campaign reality: claims, assumptions, blockers and evidence references as of a stated evidence cutoff, linked to the Review that produced it [S: HS §4, §6]. G0 materializes the Campaign Position; Market and Launch Positions become typed records at G1 [S: HS §9].

It is **not**:

- a forecast or probability that the campaign will succeed [S: HS §1.2];
- a readiness score standing in for evidence;
- causal proof: a Delta compares expected with observed change and records attribution limits [S: HS §4];
- a live record that rewrites itself or raw metrics. A newly accepted Review commits a successor version, and both remain [S: HS §5.3, §7].

Product boundary with the rest of the portfolio

Reports are rendered by the shared Report Generation Faculty, which serves every faculty [U]. Campaign execution stays in the customer's existing tools, reached only through scoped adapters [S: HS §1.2, §2.1]. Sharing with wider Faculty Systems waits for G4 and versioned conformance tests [S: HS ADR-10]. Where Hipstraw's launch testing ends and the Faculty of Launch Management begins is open (HQ7) [O].

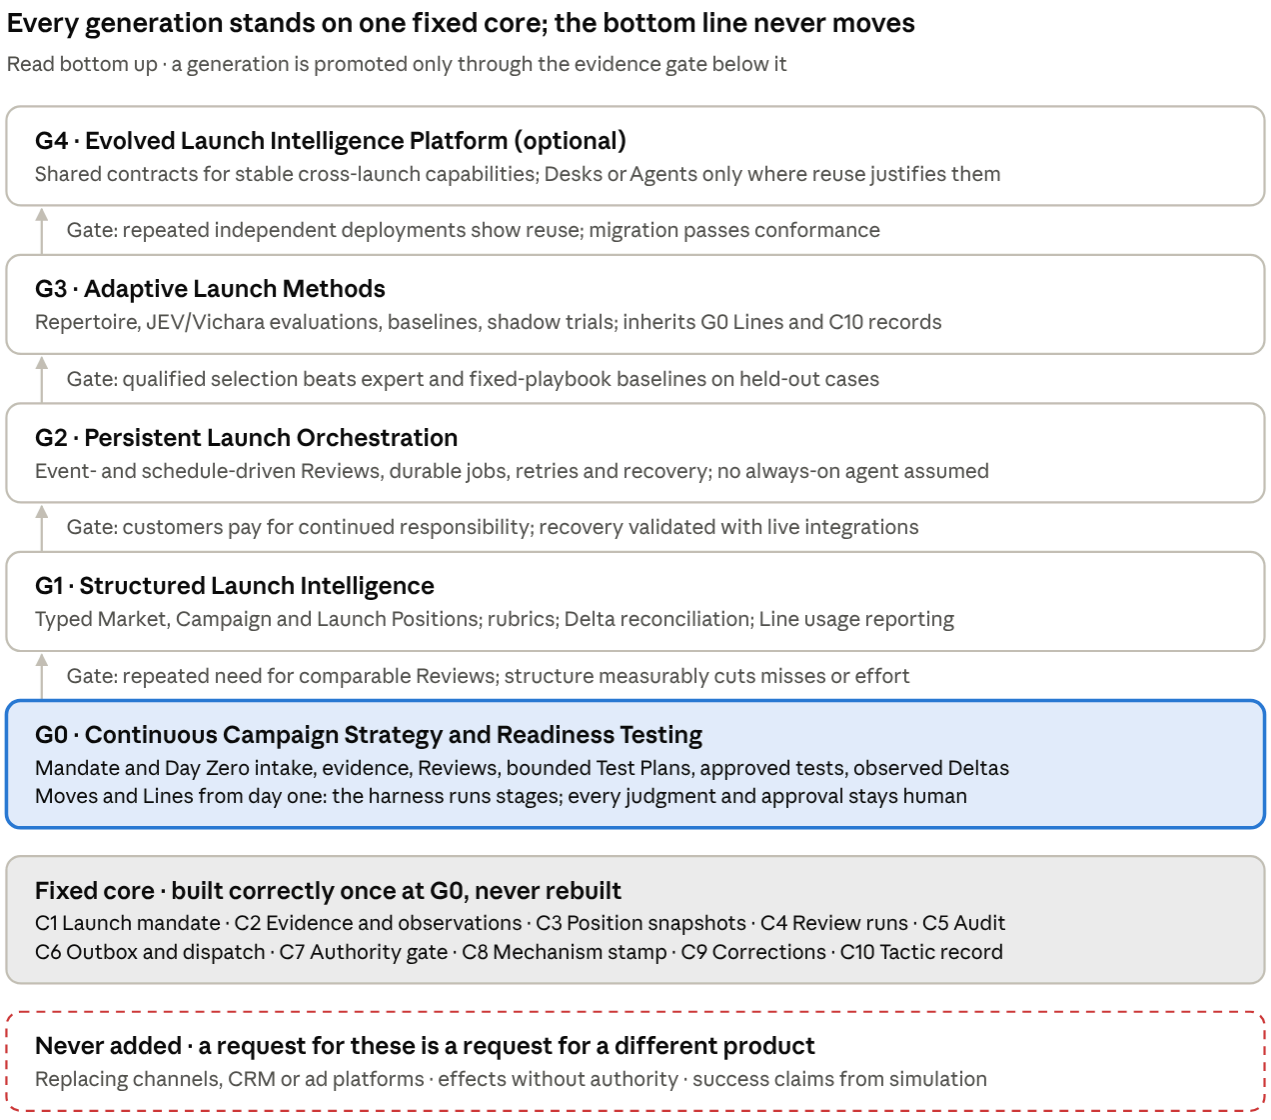
2. Design principles

HS's six principles hold [S: HS §1.3]. Four more come from the Zimplii pattern; they decide what the harness is *for* and what it builds once, before the others decide how [P].

1. **Scope is a feature.** A capability belongs only if it helps test strategy or readiness, review a Position, or coordinate an authorized test [P]. Running campaigns is declined, not deferred.
2. **Name tactics before automating them.** Moves and Lines exist as named, recorded practice from the first launch. Automating their selection or composition waits for G3 evidence [P; S: HS ADR-04, §9].
3. **Semantics before services.** Business meanings stay stable while implementations remain replaceable; a Desk is a locus of responsibility, not a deployment [S: HS §1.3, §3].
4. **Get the fixed core right once.** C1–C10 are the only parts G0 builds for the future (section 4) [P].
5. **Evidence before assertions.** Every material finding names source, recency, method and uncertainty [S: HS §1.3].
6. **Authority before effects.** Permission is validated at execution time, not only when an Action is proposed [S: HS §1.3].
7. **Determinism for invariants.** Code handles idempotency, eligibility, required fields, budgets, gates and transitions [S: HS §1.3, §8.1].
8. **LLM only where useful.** Bounded inference interprets evidence and proposes alternatives inside a declared Skill; it never decides authority, eligibility or state [S: HS §1.3, §8.1].
9. **Portable history, stamped with mechanism and tactic.** Evidence, snapshots, Reviews, approvals, Actions and observations export across rewrites, each carrying the mechanism stamp and tactic record [S: HS §1.3; P: stamp and tactic].
10. **Promote only on measured advantage.** A simpler mechanism wins unless a sophisticated one adds material value on the same cases; a generation label names what has been validated [S: HS §9.1, §10].

3. Architecture at a glance

The stack is HS's five generations standing on one fixed core, with a permanent boundary under the whole stack [S: HS §9; P: core and boundary].



Harness layers · fixed core, 5 generations, permanent exclusions

Read it from the bottom up. G0 already speaks in Moves and Lines; later generations add structure, persistence and adaptive selection, each only through its gate. The bottom line is permanent: those capabilities belong to the customer's own tools or to other products, not to a later generation of this one.

4. The fixed core

Ten elements are built correctly at G0 and never rebuilt; later generations add beside them [P]. C1–C9 give HS's record families (§6) and controls (§7, §8) a fixed shape. C10 is new: it records which Line each Test used and how the humans departed from it, so G3's

evidence base accumulates long before G3 exists [P].

#	Element	Guarantee	Later use
C1	Launch mandate	Launch identity, tenant, owner, objective, Day Zero, audiences, scope, budget, authorized channels; versioned [S: HS §6]	Every Review, gate and Delta
C2	Evidence and observation records	Immutable; source, capture and effective time, provenance, verification state, uncertainty, hash, access class; metric definition, window, denominator, confounders [S: HS §6]	Replay, Delta, attribution notes
C3	Versioned Position snapshots	Accepted snapshots immutable; successors commit only against the pinned base version [S: HS §5.3, §7]	Typed Positions at G1; Position history
C4	Review runs	Question, base Position version, evidence cutoff, findings, uncertainty, proposed delta, acceptance [S: HS §6]	Rubric comparison at G1; evaluations at G3
C5	Append-only audit and decisions	Actor, authority basis, before and after references, rationale, correlation ID [S: HS §6]	Reconstruction of consequential decisions
C6	Outbox and idempotent dispatch	State and events commit together; idempotency keys, bounded retries, reconciliation, exception queue [S: HS §7]	Persistent orchestration at G2

#	Element	Guarantee	Later use
C7	Policy and authority gate	Tenant isolation, entitlements, data rights, budgets, approvals, side-effect classes; fails closed; re-checked at dispatch [S: HS §2.1, §8]	Every E3 Move at every generation
C8	Mechanism stamp	Model family and version, prompt, rubric, connector and code versions on every Review and snapshot [S: HS §8; P: stamp]	Honest history after a mechanism is replaced
C9	Corrections and adjudication	Rejected findings and human corrections kept with reasons [S: HS §11.3]	Correction-rate promotion tests
C10	**Launch tactic record**	Per Test Plan and Action: the Line chosen (ID and version), who chose it and why, the Moves executed in order, and every deviation (skipped, added or reordered Move) with its reason	Comparing Lines (section 11); qualified Move selection at G3 [P]

C10 is cheap: one row per Test Plan plus one row per executed Move, written in the same transactions that already commit Action outcomes [P]. Without it, G3 would start with no evidence of which tactics were used or how they performed.

5. Launch tactics from day one

At G0 a **Move** is a versioned definition with eligibility, inputs, expected effects, risks and an effect class, backed by an explicit handler or template; a **Line** is a versioned definition listing Moves in ordered stages toward an intended change in Position [S: HS §4, ADR-04; P: stages]. A human chooses the Line, the harness runs its stages, and a human makes every branching judgment and approval. This gives Hipstraw its tactical language without a compiler [P].

Effect classes

Class	Meaning	Gate
E0 Internal	Deterministic check or record change; no external call	Policy gate (C7)
E1 Read external	Reads approved research sources or analytics	Connector scope in the mandate (C1)
E2 Inference	Bounded model synthesis inside a declared Skill	Versioned prompt; output is a proposal only
E3 External effect	Publishes, sends, spends, contacts or changes budgets	Scoped human approval, re-checked at dispatch; idempotency key (C6)

Typed side-effect classes and read-only-first scopes are HS's [S: HS §8, ADR-06]; the four names are [P].

Candidate Moves for the first release

Twelve candidate Moves cover the G0 flow in HS §5. They are a starting proposal; the catalogue should be drawn from the Hipstraw Market Mastery workbook (HQ1) [P].

Group	Moves	Effect classes
Evidence	M01 Capture customer asset · M02 Bounded source research · M03 Verify metric identity and freshness	E0, E1, E2
Readiness	M04 Completeness check against readiness criteria · M05 Day Zero dependency check (deadlines, required assets)	E0
Strategy	M06 Synthesize strategic assumptions and gaps · M07 Propose message or audience variants · M08 Request evidence for an unverified assumption	E2, E0
Testing	M09 Draft Test Plan (hypothesis, cohort, stop rule, budget) · M10 Dispatch approved test through an adapter · M11 Record observation and Delta	E0, E3, E1
Escalation	M12 Escalate a blocker or contested finding to the Sponsor	E0

M10 is the only E3 Move, and it runs only behind an approved Test Plan [P; S: HS §5.3].

Candidate Lines for the first release

Line	Stages (Moves)	Use when
L01 Readiness-first	M01 → M04, M05 → M12	Day Zero is near and operational blockers are the main risk
L02 Strategy validation	M02 → M06 → M07 → M09	Core campaign assumptions are unverified
L03 Test before expansion	M07 → M09 → M10 → M11	A variant needs field signal before it is scaled [S: HS App. A example]
L04 Live checkpoint	M03 → M11 → M06	At T0, T+7 or a configured event [S: HS §5.2]
L05 Contested evidence	M08 → M02 → M12	Evidence is missing or contradictory [S: HS §11.3]

Moves separated by commas run in parallel within a stage; arrows separate sequential stages [P].

How a Line runs without a compiler

1. **Scoping.** The Review question decides *what* must be tested; the Market or Campaign Manager role proposes *how* by picking a Line, or a recorded default for the launch type applies. Both are pinned on the Test Plan [P].
2. **Stages.** The harness enqueues stage 1's Moves as jobs. When every job in a stage is terminal, the next stage starts.
3. **Judgment.** Between stages a human may skip a Move, add one from the catalogue, or stop the Line early. Each choice is a deviation in the tactic record (C10).
4. **Gates.** An E3 Move waits for Test Plan approval; the gate re-checks budget, channel scope and approval at dispatch [S: HS §5.3, §7].
5. **Exit.** Every Line ends in a Review that proposes a successor Position. Only acceptance commits it; no Line commits a Position or releases a report.

Checking Line definitions

Line definitions live in the repository and are checked in CI by a short list of rules [P]: every Move ID exists and is versioned; every E3 Move sits behind a Test Plan approval step and an authorized mandate channel; every Line ends in a Review; no E2 output becomes a Position claim without evidence references. This is a lint, not the compiler: no intermediate representation, branches or qualification ladder.

6. Domain model

The grammar is HS's (§4): evidence → proposed Position → Review → decision → selected Line and Move → authorized Action or test → observed evidence and Delta → next Position [S: HS §4]. This architecture adds one thing: tactics are first-class records from G0, pinned on every Test Plan [P].

Core relationships

- A **Tenant** has **Launches**; each Launch carries a **Mandate**: owner, Day Zero, objective, scope, budget and authorized channels [S: HS §6].
- A Launch has **Positions** (Campaign at G0; Market and Launch typed at G1), each a chain of versioned **Snapshots** [S: HS §9].
- A **Review** pins a base Position version and an evidence cutoff, reads **Evidence**, and proposes **Findings** and a successor Snapshot [S: HS §7].
- A **Decision** accepts or rejects the Review. A selected **Line** of **Moves** becomes a **Test Plan**, whose approval permits **Actions** [S: HS §4, §5.3].
- **Observations** record what happened; a **Delta** compares expected with observed change and carries attribution notes [S: HS §6].
- Every Test Plan has a **Tactic Record** of its Line, executed Moves and deviations [P].

Entities

Entity	Meaning	Introduced
Tenant, Principal, Launch, Mandate	Customer, person or agent acting for it, the bounded commercial undertaking and what it permits	G0
Evidence Record, Observation	Sourced claims, metrics and events with lineage and uncertainty (C2)	G0
Position Snapshot (Campaign)	Accepted, versioned representation of campaign reality (C3)	G0
Review Run, Finding, Correction	Assessment, its findings and human adjudication (C4, C9)	G0
Decision, Audit Event	Acceptance, authorization and history (C5)	G0
Play / Playbook	Business-level plan or reusable pattern; template only, no compiler [S: HS §4]	G0

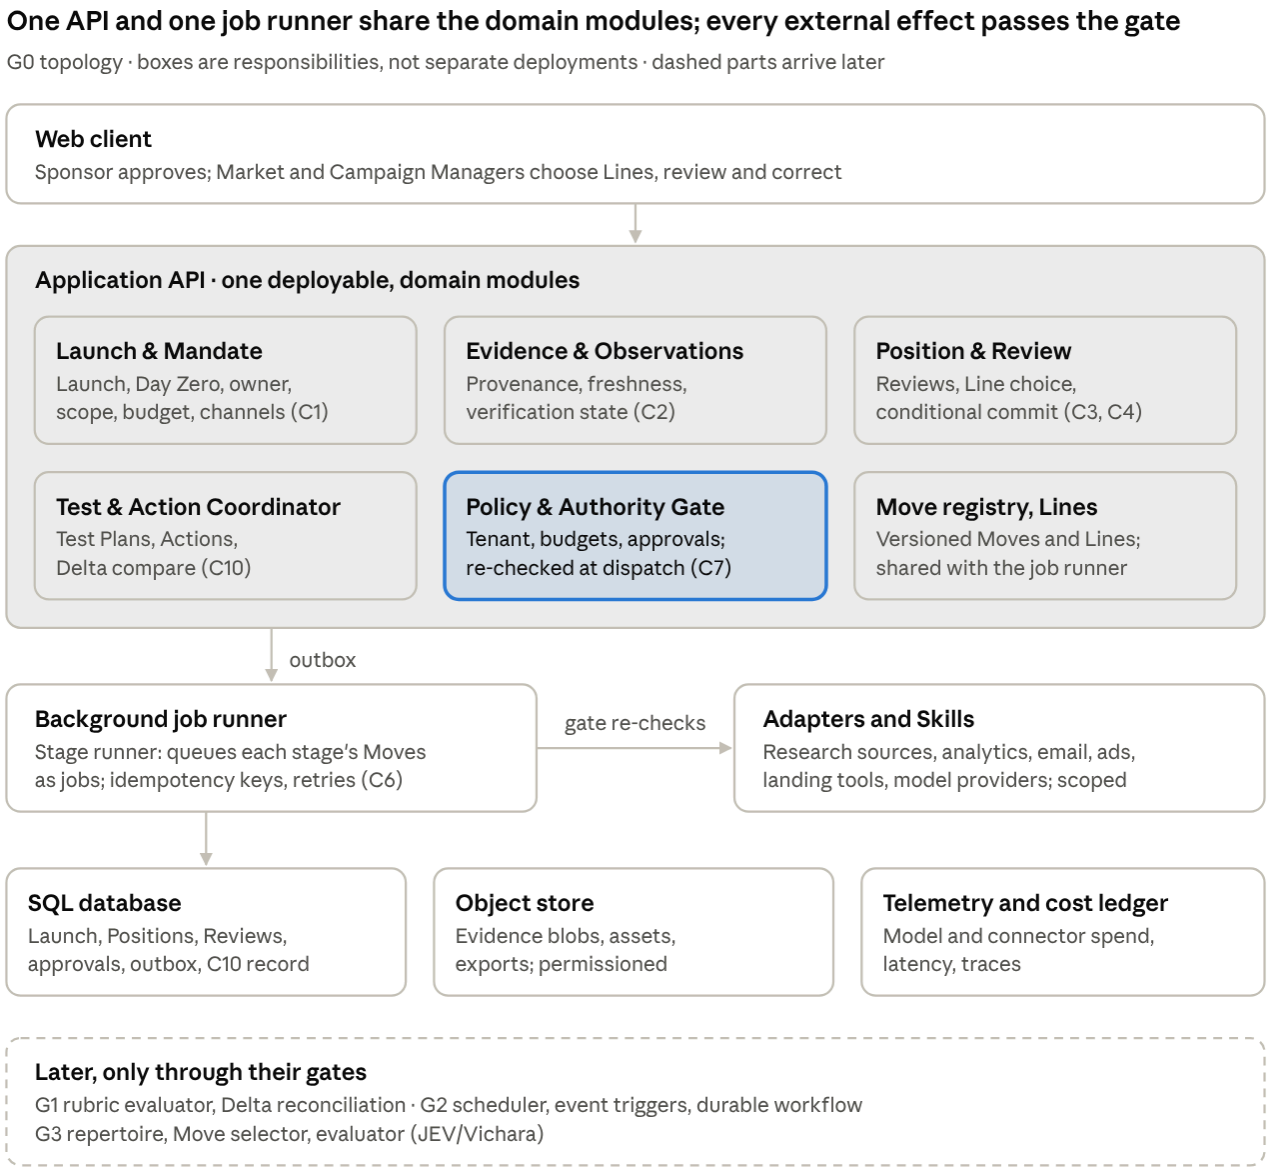
Entity	Meaning	Introduced
Move Definition	Versioned Move: ID, group, effect class, eligibility, inputs, expected effects, risks	G0 [P]
Line Definition	Versioned Line: ordered stages of Move IDs, use-when and exclude-when conditions	G0 [P]
Test Plan, Action	Bounded test and its concrete executable instances	G0
Tactic Record, **Move Execution**, **Deviation**	C10: what was chosen, what ran, what humans changed and why	G0 [P]
Delta	Expected versus observed change; never causal proof	G0
Report	Rendered by the shared Report Generation Faculty from an accepted Snapshot [U]	G0
Usage Record	Model, connector and job cost per Review and Test [S: HS §12]	G0
Typed Market, Campaign and Launch Positions; Rubric; Delta reconciliation	Structured launch intelligence [S: HS §9]	G1
Schedule, Trigger, Durable Job, Dependency Watch	Persistent orchestration [S: HS §9]	G2
Candidate Repertoire, Evaluation, Baseline, Qualification Record	Adaptive methods, including JEV/Vichara evaluations [S: HS §9]	G3
Shared contract package	Extracted cross-launch capability [S: HS §9]	G4

An accepted Snapshot is never edited, and a Delta never alters a Snapshot or a raw Observation [S: HS §5.3, §7].

7. Components and where they run

The deployment shape is HS's G0 topology: one authenticated web client, one application API with Launch, Evidence, Review, Test and Policy modules, one SQL database, a permissioned object store, one managed background-job mechanism, selected

connectors and an observability and cost ledger [S: HS §11.1]. This architecture adds Line choice, a stage runner, Line lint and the tactic record to G0 [P].



Full G0 architecture · client, API modules, job runner, adapters, storage, later parts

The web client and the background runner call the same domain modules [S: HS §2.2]. The stage runner queues each stage's Moves as jobs and waits for them to finish before the next stage; the tactic record sits in SQL beside the other core records. G2 parts arrive only after the G1 → G2 gate, and G3 parts with adaptive methods.

Components new or changed in this architecture

Component	Responsibility	Runs in	Generation
Line choice	Offers Lines that fit the Review question and launch type; records who chose and why	API, Position & Review module	0
Stage runner	Queues each stage's Moves as jobs; advances when a stage is terminal; records deviations	Background job runner, as an outbox consumer	0
Move registry and Line definitions	Versioned Move handlers and templates; versioned Line definitions	Shared modules	0
Line lint	Checks Line definitions against effect-class and gate rules	CI	0
Position commit	Conditional commit against the pinned base version; rejects stale proposals [S: HS §7]	API	0
Delta compare	Expected versus observed change with attribution notes	API, Test & Action module	0
Rubric evaluator, Delta reconciliation	Deterministic rubrics and typed Position comparisons	API	1 [S: HS §9]
Scheduler, event triggers, durable workflow	Event- and schedule-driven Reviews, multi-day waits, recovery	Worker or managed engine	2 [S: HS §10]
Repertoire, Move selector, evaluator	Qualified Move and Line selection by applicability	Shared modules, worker, lab	3 [S: HS §9]

All other components (Launch & Mandate, Evidence & Observations, Test & Action Coordinator, Policy & Authority Gate, adapters and Skills, persistence and telemetry) keep HS's responsibilities and prohibitions [S: HS §2.1].

8. Launch lifecycle and operational behaviour

HS's state machines stand [S: HS §5.3]. This architecture adds two things inside them: Line selection when a Review leads to a Test Plan, and stage-by-stage progression while the Test Plan runs [P]. Recovery mechanics (idempotency keys, reconciliation, outbox, exception queue) are unchanged [S: HS §7].

State machines

Object	States	Hard invariant
ReviewRun	requested → gathering → drafted → validated → accepted / rejected / failed	Only accepted results may update an approved Position [S]
PositionSnapshot	proposed → accepted → superseded	Accepted snapshot immutable; replacement creates a new version [S]
TestPlan	draft → eligible → awaiting_approval → approved / rejected → running → evaluated / failed / cancelled	Eligibility is not authorization; approval is scope-bound [S]. A Line is pinned at *draft* [P]
Action	proposed → eligible → authorized → queued → running → succeeded / failed / cancelled	No external effect without current permission and an idempotency key [S]
TacticRecord	opened → staged (per stage) → completed / abandoned	Every Test Plan and Action has one; every executed job maps to a Move ID [P]

End-to-end flow

1. **Mandate.** Create the Launch with tenant, owner, objective, Day Zero, campaign strategy and mandate limits; record input and schema versions [S: HS §5.1].
2. **Evidence.** Ingest customer assets and selected third-party data; keep provided claims, verified evidence and unknowns apart [S].
3. **First Position.** Build the proposed Campaign Position; run deterministic completeness and readiness checks; run bounded synthesis only if justified [S].
4. **Review.** Findings separate strategy validity from operational readiness, with evidence, uncertainty, blockers and open assumptions [S].

5. **Line and Test Plan.** The Manager role picks a Line; its stages draft a limited Test Plan with hypothesis, cohorts, method, success and stop criteria, budget, risks and approvals [S; P: Line].
6. **Approval.** The Sponsor accepts or rejects findings and authorizes the test; the gate checks the current mandate and connector scope [S].
7. **Execution and Delta.** Stages run; attempts and observations are captured; expected and observed Delta are compared; deviations land in C10 [S; P: C10].
8. **Commit.** The accepted successor Position and decision are committed: continue, change, delay, re-test or stop. Original evidence and corrections are kept [S].

Live checkpoints. At T0, T+7 or a configured event, a Review obtains permitted measurements, verifies event identity, timestamps and freshness, compares actuals with the baseline, and separates delivery defects from strategy signals [S: HS §5.2]. Any new intervention passes the same eligibility, approval and dispatch gates as a prelaunch test.

Failures

HS's operational scenarios apply: contested evidence returns *unknown* with an evidence request, never false confidence; duplicates deduplicate; stale commits are re-reviewed; revoked approvals block dispatch; uncertain timeouts are reconciled, not blindly repeated [S: HS §11.3]. Two failures are new [P]:

Failure	Handling
Chosen Line references a Move whose handler version is withdrawn	Selection refused; the Manager picks another Line or runs Moves individually, recorded as deviations
Humans deviate so far that no Line stage remains	Test continues as a human-directed Test; the tactic record marks the Line abandoned

9. Security, privacy and tenancy

HS's controls apply in full: tenant scope on every API, query, artifact, job, index and model or tool call; recommendation separated from authorization; retrieved pages, assets and tool outputs treated as untrusted data; evidence labelled unverified, reported, observed or primary-source verified; and every model call logged with version, prompt, inputs, tokens, spend and corrections [S: HS §8]. Read-only sources come first, and write scopes are approved separately [S: HS ADR-06].

Data classes in the first release

Class	In first release?	Handling
D1 Public and third-party research sources	Yes	Untrusted input; to a model only under the approved provider policy [O]
D2 Customer strategy and assets	Yes	Scoped object-store access; field-level access and retention [S: HS §8]
D3 Campaign metrics and analytics observations	Yes	Source ownership, metric definition and capture time preserved; raw values never rewritten [S: HS §7]
D4 Personal data (audience and contact lists)	**Only** inside an E3 test with lawful data rights	Minimized; masked in logs and telemetry; no secondary training use without express rights [S: HS §8; P]
D5 Secrets and connector credentials	Yes	Secret manager only; scoped per adapter [P]

Roles

Capability	Sponsor / Approver (GTM leader)	Market Manager (agent role)	Campaign Manager	Domain owner	Operator
Approve objectives, strategy changes, Test Plans and E3 Actions	Yes	No	No	No	No
Prepare and run the Market Review; propose a Market Position and a Line	No	Yes	No	No	No
Maintain Campaign Position and readiness Review; coordinate approved tests	No	No	Yes	No	No
Change Move or Line definitions, rubrics and rules	No	No	No	Yes, by reviewed change	No

Capability	Sponsor / Approver (GTM leader)	Market Manager (agent role)	Campaign Manager	Domain owner	Operator
Retry dead-lettered jobs; see metadata only	No	No	No	No	Yes

The Sponsor and Market Manager roles are [U]; the Campaign Manager role is [S: HS §3]; the others are [P]. Agent roles never hold external-effect authority, and their tools are scoped and mediated [S: HS §3]. Whether any material Position change may be accepted without a human is ADR-02 [O].

Data rights and scopes

1. Which sources and connectors may be read (E1).
2. Which channels may be written to and which budgets spent (E3), each with expiry and revocation [S: HS §8].
3. Whether the launch may be used, de-identified, in cross-launch evaluation. Off unless the customer grants it [S: HS §8; P: default]. This scope decides how fast Lines can later be compared (section 11).

10. Evolution by generation

The five generations are HS's, each promoted only on its evidence gate [S: HS §9, §9.1]. Generations are acceptance levels, not exclusive work streams: a G0 team may research G3 methods in a lab, but the production label names only what has been validated [S: HS §9.1]. This architecture adds one inheritance: G0's Line definitions and tactic records become G3's inputs [P].

G0 — Continuous Campaign Strategy & Readiness Testing

Aspect	Design
Builds	Sections 4–9: fixed core C1–C10, candidate Moves and Lines run as ordered stages, human acceptance of every material Position change, approved tests and stored outcomes
Excludes	Proprietary compiler, persistent agent sessions, federation [S: HS §9]

Aspect	Design
Exit test (G0 → G1)	Customers repeatedly need comparable Reviews and evolving launch status; structured representation measurably reduces misses, corrections or effort [S: HS §9.1]

G1 — Structured Launch Intelligence

Aspect	Design
Adds	Typed evidence; versioned Market, Campaign and Launch Positions; standardized Review contracts; deterministic rubrics; Delta reconciliation [S: HS §9]. **Line usage and deviation reporting** from C10; a recorded default Line per launch type [P]
May replace	Free-form Review judgment with rubric-guided judgment, after a comparison on shared cases (section 11)
Exit test (G1 → G2)	Customers pay for continued responsibility; triggering, approval, idempotency and recovery validated with live integrations [S: HS §9.1]

G2 — Persistent Launch Orchestration

Aspect	Design
Adds	Event- and schedule-driven Reviews, durable jobs, dependency monitoring, authorized Lines and Moves, retries, recovery and managed continuity [S: HS §9]
Does not assume	An always-on LLM agent [S: HS §9]
Mechanism choice	Graph workflow, event-driven controllers or a durable workflow engine, chosen by the comparative test in section 11 [S: HS §10]
Exit test (G2 → G3)	Qualified selection beats expert or fixed-playbook baselines on held-out or prospective cases within explicit applicability bounds [S: HS §9.1]

G3 — Adaptive Launch Methods

Aspect	Design
Adds	Candidate repertoire, JEV/Vichara evaluations, baselines, simulated and shadow trials, qualified Move selection by applicability [S: HS §9]
Inherits	G0 Line definitions become selection and compiler inputs: stages become graph segments, and deviations recorded in C10 show where branches are needed. Tactic records become the evidence base [P]
Never	Claims field effectiveness from generation or simulation alone [S: HS §9]
Exit test (G3 → G4)	Repeated independent deployments show reuse, lower total ownership cost or stronger outcomes, and migration passes conformance [S: HS §9.1]

G4 — Evolved Launch Intelligence Platform (optional)

Stable cross-launch capabilities become shared contracts, with specialized Desks or Agents only where reuse justifies them [S: HS §9]. There is no mandatory cross-product federation or full Faculty ontology [S: HS §9].

Never added to this product

Replacing campaign channels, analytics, CRM, email or ad platforms; external effects without specific authority; and success claims drawn from generation, simulation or engagement metrics alone [S: HS §1.2, §9; P: permanence]. A request for any of them is a request for a different product.

11. Comparing tactics and replacing mechanisms without a compiler

Before G3, Lines are compared by observing how they perform in recorded launches. That improves Line definitions but cannot prove one Line is better, because Managers choose Lines for particular kinds of launch. Proof needs the controlled, preregistered comparisons that arrive with G3 [P; S: HS §9, §12].

What the tactic record makes measurable (G1)

Measure per Line	What it suggests
Blockers found before Day Zero, and blockers missed until later	Whether the Line surfaces readiness problems in time [S: HS §12]
Hypotheses resolved and time-to-signal	Whether the Line's tests produce decisions quickly [S: HS §12]
Decision changes (continue, change, delay, re-test, stop)	Whether the Line moves the launch, not just the record
Deviation rate and most common deviations	Whether the Line matches how Managers actually work
Human minutes, corrections and cost per accepted Review or Test	Effort and quality [S: HS §12]

A high deviation rate on one stage is the clearest signal: the Line is wrong for that kind of launch. The domain owner revises it as a new version through reviewed change; old Test Plans keep the version they used [P].

Limits of observational comparison

- **Selection bias.** L01 is chosen when Day Zero is near, so it cannot fairly be compared with L02 across all launches. Compare within a launch type and record every choice's reason (C10).
- **Weak causal attribution.** Campaign outcomes carry confounders; measure correctness and operational quality separately from campaign effectiveness [S: HS §12].
- **Replay can't run tests that never ran.** Replay re-tests Review judgment on stored evidence, not a Line that would have run a different field test.
- **Small samples.** Only launches with scope-3 data rights can be used across customers; early counts will be small.

These limits are why controlled comparison, with baselines, held-out or prospective cases and applicability bounds, belongs to G3 [S: HS §9.1].

Replacing a mechanism

A candidate mechanism is preregistered with its question, baseline, shared case set, evidence cutoff, success metric, guardrails, budget, stopping rule and reviewer protocol [S: HS §12]. It runs the same authorized cases as the incumbent and is measured on precision, missed blockers, correction rate, latency, total cost and recoverability; a simpler

mechanism wins if the sophisticated one adds no material value [S: HS §10]. Old Positions keep their original mechanism stamp (C8). A revised Line definition is a new version, not a mechanism replacement, and needs only reviewed change [P].

12. Verification and observability

HS's operational acceptance scenarios and contract-first fixtures apply unchanged [S: HS §11.3, App. D]. This architecture adds tests that protect the product's scope and its tactic record, and measures that show how Lines perform [P].

Added acceptance fixtures

Fixture	Expected result
Manager skips a Move in L01	Deviation recorded with reason; Review still proposes a successor; stamp lists the deviation
Line places an E3 Move without an approval step	Line lint fails in CI
Test Plan reaches M10 after approval is revoked	Dispatch blocked; deviation and reason recorded; no external effect
Customer asks for a success probability	No such field exists in the Position or report schema
Data-rights scope 2 grants no channel	M10 cannot run; Lines that require it are not offered

Added architecture checks

- The Position and report schemas contain no success-probability or forecast fields [P].
- No adapter holds a write scope without a matching mandate entry and an approval reference [P; S: HS §8].
- Every Test Plan row has a tactic record row; every executed job maps to a Move ID; every Review row has a mechanism stamp [P].

Measures

HS's measurement categories stay: judgment quality, test quality, operational quality, economics, outcome impact with explicit confounders, and pass/fail safety gates [S: HS §12]. This architecture adds the per-Line measures from section 11 and **scope-pressure**

tracking: requests for out-of-scope capabilities (autonomous publishing, budget control, success forecasts, channel replacement), counted by type [P]. That count is evidence for future scope decisions, not a backlog.

13. Decision register

Twelve decisions. HS's ADR-01 to ADR-10 stay open unless a decision below answers them: HR05 proposes ADR-04's answer, HR12 adopts ADR-06's default, and HR09 narrows ADR-02 [P].

ID	Decision	Section	Approver	Status
HR01	The G0 product is Continuous Campaign Strategy & Readiness Testing, defined by the scope statement	1	Product owner	Needs approval
HR02	Replacing channels, analytics, CRM, email or ad platforms, and external effects without specific authority, are permanently outside this product	1, 10	Product owner	Needs approval
HR03	A Position is an accepted, versioned, evidence-linked snapshot; never a success forecast; a Delta is never causal proof	1	Product + domain owner	Proposed
HR04	Build the fixed core C1–C10 at G0	4	Architecture owner	Proposed
HR05	Moves are versioned definitions with effect classes E0–E3 and Lines are versioned stage definitions from G0; a human chooses; Line lint in CI (ADR-04)	5	Architecture + domain owner	Proposed
HR06	Compiler and adaptive selection stay at G3 behind the G2 → G3 gate	10	Architecture owner	Proposed
HR07	Add the launch tactic record (C10) to the fixed core	4	Architecture owner	Proposed

ID	Decision	Section	Approver	Status
HR08	First-release repertoire: the 12 candidate Moves and 5 candidate Lines in section 5, pending the workbook (HQ1)	5	Domain owner	Open
HR09	The GTM leader is Sponsor/Approver; a Market Manager agent prepares and runs the Market Review for the Sponsor's decision (narrows ADR-02)	9	Product owner	Decided
HR10	Reports are rendered by the shared Report Generation Faculty	1	Product owner	Decided
HR11	Every Review and report carries the scope statement	1	Product owner	Proposed
HR12	Read-only adapters first; each write scope approved separately (ADR-06)	9	Product + security owner	Proposed

14. Open questions and dependencies

HS's ADR-01 to ADR-10 remain open except where section 13 answers them. This architecture adds questions about tactics, Position types and the faculty boundary. Implementation design can start once HQ1–HQ4, ADR-01 and ADR-06 are answered [P].

#	Question	Needed for	Owner
HQ1	Which Moves and Lines from the Hipstraw Market Mastery workbook form the first-release repertoire, and which Line is the default per launch type?	HR08; Line choice	Domain owner
HQ2	Which deviations must a Manager explain, and with what reason codes?	C10	Domain owner

#	Question	Needed for	Owner
HQ3	Does G0 materialize only the Campaign Position, or the Market Position too, given the Market Review in HR09?	Section 6; HR09	Product owner
HQ4	Which live checkpoints are default: T0, T+7, or customer-configured events only?	L04; section 8	Product owner
HQ5	Exact wording of the scope statement and of what a Position is not	HR11; report template	Product owner + legal review [A]
HQ6	What measured evidence clears the G0 → G1 gate: which misses, corrections or effort reductions, and how many launches?	Section 10	Product owner
HQ7	Where does Hipstraw's launch testing end and the Faculty of Launch Management begin?	Section 1 boundary	Product + architecture owner

Dependencies before implementation design

- [ ] Approve or amend HR01–HR07, HR11 and HR12.
- [ ] Answer HQ1–HQ4, ADR-01 and ADR-06.
- [ ] Write the Line definitions and Move handler contracts for the approved repertoire.
- [ ] Extend HS Appendix D fixtures with section 12's additions.
- [ ] Define the portable record envelope, including C10, from HS Appendix A (ADR-01).
- [ ] Freeze the G0 vertical-slice contract: Review → Test → Action and Observation → Position update [S: HS §11.2, §14].